

CS 423 — Computer Architecture

Bilkent University — Fall 2026

Homework 1: In-Order Pipeline

100 points • Out: September 24, 2026

Due: October 18, 2026, 23:59 (Ankara time)

Learning Objectives

After completing this homework, you will be able to:

- run `gem5` syscall-emulation (SE-mode) simulations in the pinned Docker environment and interpret `stats.txt` (instructions, cycles, CPI/IPC, simulated vs. host time);
- explain what `TimingSimpleCPU` (a non-pipelined CPU) and `MinorCPU` (an in-order pipelined CPU) each model, and compare the simulated non-pipelined CPU with an idealized single-cycle design;
- quantify the throughput benefit of pipelining and reason about the clock-period side of the trade-off;
- provoke, observe, and explain data hazards (especially load-use hazards) and control hazards with microbenchmarks you write yourself;
- present measured results as clear plots with mechanism-level analysis.

GenAI Policy

GenAI tools are **allowed** as helpers for coding, for preparing data visualization scripts, and for proofreading your final text. They are **not allowed** to take over any task completely (e.g., implementing mechanisms, conducting simulations, analyzing data, or writing critical reviews). You are the sole author of everything you submit. In HW1 specifically, a tool may remind you of GCC inline-assembly syntax, but the *hazard structure* of your Task 4 microbenchmarks and every prediction and explanation in your report must be your own. Your report must include a short GenAI disclosure statement (see the report template).

Environment Setup

All experiments run in the course `gem5` Docker environment: `gem5 v25.1.0.1`, RISC-V, with the cross toolchain preinstalled (see `hw1/README.md` for details and for a build-it-yourself fallback):

```
git clone https://github.com/SAL-Research/cs423-fall2026-homeworks.git
docker pull s4lbot/cs423-gem5:v25.1.0.1
cd cs423-fall2026-homeworks/hw1
docker run --rm -u "${id -u}:${id -g}" -v "$PWD":/hw -w /hw \
  s4lbot/cs423-gem5:v25.1.0.1 bash scripts/smoke_test.sh
```

The homework repository is *mounted* into the container; the image itself contains only the environment (`gem5` and the toolchain). Therefore, `git pull` always gets you the latest assignment files without re-downloading the image. Results obtained with other `gem5` versions may differ and will not be regraded.

Background

A *single-cycle* processor completes every instruction in one clock cycle, so its clock period must cover the longest path any instruction takes through the datapath. A *pipelined* processor divides the datapath into stages separated by pipeline registers and works on a different instruction in each stage. Its clock period only has to cover the slowest stage (plus the pipeline-register overhead), and ideally one instruction still completes every cycle. The *equation of the execution time* combines the instruction count, the CPI, and the clock period:

$$\text{execution time} = \text{IC} \times \text{CPI} \times T_{\text{clk}},$$

where IC is the number of executed instructions, CPI is the average number of cycles per instruction, and T_{clk} is the clock period.

Overlapping instructions creates *hazards*. A *data hazard* occurs when an instruction needs a result that an earlier instruction has not produced yet. *Forwarding* (bypassing) sends a result directly to the dependent instruction instead of waiting until it is written to the register file, but it cannot deliver a loaded value before the memory access has returned it: an instruction that uses a load's result right away may still have to wait (a *load-use* stall). A *control hazard* occurs because the next instruction to fetch after a branch is not known until the branch is resolved. A *branch predictor* lets the pipeline keep fetching along a guessed path; when the guess is wrong, the wrongly fetched instructions are discarded, and the lost cycles are the *misprediction penalty*. In this homework, you measure these effects on a cycle-level simulator.

You will compare two gem5 CPU models on identical systems (same clock, same small L1 caches, same DDR3 memory):

Role	gem5 model	What it models
non-pipelined baseline	TimingSimpleCPU	one instruction at a time, no overlap
in-order pipeline	MinorCPU	four-stage in-order pipeline (Fetch1, Fetch2, Decode, Execute) with forwarding and a branch predictor, configured as a <i>scalar</i> pipeline (at most one instruction per stage per cycle)

By default, gem5's MinorCPU is two instructions wide. The provided `configs/pipeline.py` sets it to one instruction per cycle so that it matches the scalar pipeline described above. Keep that setting for every number you report, except in the parameter probe of Task 3.3.

Task 1: Environment bring-up and first simulation [10 points]

Set up the environment (see the setup box above) and run the smoke test. It cross-compiled the provided workloads and runs `hello` on the baseline CPU; the results are written to `results/smoke/stats.txt`.

From that run, report and answer the following (Q1.1–Q1.4):

Q1.1 `simInsts`: the number of committed instructions.

Q1.2 `system.cpu.numCycles`: the number of simulated CPU cycles.

Q1.3 `simSeconds` vs. `hostSeconds`: what does each measure, and why are they so different?

Q1.4 Why is `simInsts` so much larger than the handful of instructions visible in `hello.c`? (Hint: the workloads are statically linked.)

Task 2: Characterize the non-pipelined baseline [20 points]

Using the *provided, complete* `configs/baseline.py` (TimingSimpleCPU at 1 GHz, 16 KiB 2-way L1 instruction and data caches, DDR3-1600 memory), run the three provided workloads and extract

the instruction count, cycle count, CPI, and simulated time of each. Use `scripts/run_all.sh`, which also calls `scripts/extract_stats.py`. The pipeline runs in that script will fail until you complete Task 3; this is expected.

- **compute**: an ALU-bound loop with four independent accumulators;
- **memstream**: a walk over an 8 MiB array with a 64-byte stride (one access per cache block), far larger than the L1 data cache;
- **branchy**: a loop whose branch direction depends on pseudo-random data.

Answer the following (**Q2.1–Q2.3**):

- Q2.1** Rank the three workloads by baseline CPI. This CPU executes one instruction at a time, so none of its cycles are caused by pipeline hazards. Where do the cycles go in each workload? Explain the ranking from the workloads' source code.
- Q2.2** How would a *single-cycle* processor (see the Background) execute these workloads, and what would that imply for its clock period?
- Q2.3** Support your explanation in Q2.1 quantitatively with at least one statistic from `stats.txt` other than `simInsts` and `numCycles` (for example, from the L1 data-cache statistics under `system.cpu.dcache`).

Task 3: Bring up the in-order pipeline and race it [25 points]

3.1 Complete configs/pipeline.py (8 pts)

The skeleton provides everything except the CPU: instantiate `RiscvMinorCPU`, attach the provided L1 caches, and create the interrupt controller (the `TODO(student)` block walks you through these steps; `baseline.py` is your reference). Keep every non-CPU parameter identical to the baseline, and keep the provided scalar configuration. Then run all three workloads on both CPUs (`scripts/run_all.sh`) and report, for each workload, the CPI and IPC of both CPUs and the pipeline's CPI normalized to the baseline's CPI (plots: `scripts/plot_results.py`). Both CPUs run at the same 1 GHz clock, so this ratio compares cycles per instruction only; it is *not* the speedup of a pipelined design over a single-cycle design, which also depends on the clock period (Task 3.2). Explain each workload's result at the level of mechanisms.

3.2 The clock-period side of the bargain (7 pts)

The simulations compared the two CPUs *at the same clock frequency*, which ignores the main benefit of pipelining: a shorter clock period. Assume a classic five-stage scalar RISC pipeline (instruction fetch, decode and register read, execute, memory access, write-back) built from components with the following delays. In the single-cycle design, an instruction's delay is the sum of the delays of the components it uses.

Component	Delay
Instruction fetch (memory)	200 ps
Register read	100 ps
ALU operation	200 ps
Data access (memory)	200 ps
Register write	100 ps

- Q3.1** What clock period must the single-cycle implementation use, and what clock period can the five-stage pipelined implementation use? (Either ignore the pipeline-register overhead or add an overhead of your choice; in both cases, state your assumption explicitly.)
- Q3.2** Apply the equation of the execution time. Using $\text{CPI} = 1$ for the single-cycle machine and your *measured* `MinorCPU` CPI for the pipelined machine, compute the execution-time ratio of the two machines for each workload. Which machine wins on which workload, and by how much?

Q3.3 Your measured CPI was collected at 1 GHz, with realistic memory latency, on `MinorCPU`'s four-stage pipeline. Explain why plugging it into the comparison of Q3.2 is only an approximation, and in which direction it biases the result for `memstream`.

3.3 Parameter probe (5 pts)

Change *one* aspect of the `MinorCPU` pipeline with the `--minor-param NAME=VALUE` option of `pipeline.py`; repeat the option to change several parameters that belong to the same aspect. Add the run to `scripts/run_all.sh` (a commented template line is provided there). All parameters and their defaults are listed in `/opt/gem5/src/cpu/minor/BaseMinorCPU.py` inside the container. Suggestions:

- pipeline width: set `decodeInputWidth`, `executeInputWidth`, `executeIssueLimit`, and `executeCommitLimit` to 2 (gem5's default two-wide configuration);
- `executeBranchDelay` (default 1): the delay between Execute deciding to redirect fetch after a branch and Fetch1 reacting;
- the latency between two adjacent stages, e.g., `decodeToExecuteForwardDelay` (default 1 cycle).

Report the CPI with and without your change on at least one workload, and explain the mechanism behind the difference.

3.4 Find one stall in the pipeline trace (5 pts)

Capture a short `MinorTrace` window with `scripts/trace_pipeline.sh` (its header explains how to read the trace). Locate one cycle range in which the pipeline stalls, quote those trace lines in your report (a few lines, not pages), and explain what the pipeline was waiting for.

Task 4: Hazard microbenchmarks — provoke stalls on purpose [25 points]

This is the only task in which you write a program from scratch. The templates in `workloads/task4/` compile and run as shipped; you replace their placeholder kernels. For *each* experiment below, **write down your predicted outcome before running it**, then measure, then explain. The prediction, the measurement, and the mechanism carry the points; the elegance of your assembly code carries none.

4.1 Load-use hazards and dependency chains (13 pts)

In `loaduse.c`, implement two kernels with the *same* instruction count and mix (8 loads and 8 adds per iteration, in two groups of 4 loads followed by 4 adds): a *dependent* kernel whose loads form a pointer chase, in which each load's result is the address that the next load dereferences (`ld p, 0(p)`; the scaffold pre-builds the ring of pointers), and an *independent* kernel that loads from the same array at fixed offsets. The 4 KiB array fits in the L1 data cache, so what you measure is pipeline behavior, not cache misses. Measure the CPI of both kernels on *both* CPUs (4 runs; `run_all.sh` covers them), and explain:

- How does the dependency structure affect each CPU, and why?
- Why can forwarding not remove the wait for a load's data?
- From your measured extra cost per load, how much of the load-to-use latency does the in-order pipeline hide in the independent kernel?

For the last question, first determine the effective load-to-use latency of an L1 hit on `MinorCPU` and explain how you derived it. `baseline.py` sets the L1 `tag_latency`, `data_latency`, and `response_latency` to 1 cycle each; which of these apply to a hit is decided by the gem5 cache model (see the `sequential_access` parameter in `/opt/gem5/src/mem/cache/Cache.py`).

4.2 Predictable vs. unpredictable branches (12 pts)

In `brpair.c`, implement the body of `run_loop()`: a conditional branch on `cond[i]` with cheap work on each of its two paths. The same loop runs over one of two precomputed condition arrays, selected on the command line: a repeating pattern (`pred`) or pseudo-random bits (`rand`). Measure on `MinorCPU`: the CPI and the two branch-predictor statistics `condPredicted` and `condIncorrect` (both under `system.cpu.branchPred` in `stats.txt`). When you need the number of branches per iteration, count them in the kernel's disassembly (`riscv64-linux-gnu-objdump -d`) rather than deriving it from `condPredicted`: that counter also includes the branches of the array-fill loop and branches fetched on mispredicted paths.

Compute the misprediction rate of both variants and estimate the average penalty (in cycles) per misprediction. Base the estimate on the difference in *total cycles* between the two variants divided by the difference in `condIncorrect`. The two variants generally do not commit the same number of instructions (the two paths of your branch can differ in length, and the variants take them in different proportions), so a CPI-based estimate would silently mix that effect in; state how you account for it. What could the predictor learn, what could it not learn, and why? (Treat `MinorCPU`'s default branch predictor, which is enabled out of the box, as a black box.)

Task 5: Report [20 points]

Submit one PDF that follows `report/report-template.md` and contains all plots (generated with the provided script or your own), the answers to every question above, a comparison of the simulated non-pipelined CPU (`TimingSimpleCPU`) with the single-cycle design of Task 3.2 (how do they differ, and does the difference affect your conclusions?), and a reproducibility statement. Task 5 points reward the quality of interpretation and presentation; the raw numbers are graded in Tasks 1–4.

Deliverables and Submission

Submit one archive, `hw1_<studentID>.zip`, via **Moodle**. It must contain your complete `hw1/` directory, so that the graders can run `reproduce.sh` from it directly (run `make -C workloads clean` first to leave out the compiled binaries). In particular, it must contain:

- **report.pdf**: your report (Task 5), including the answers to all questions;
- **configs/**: including your completed `pipeline.py`;
- **workloads/**: including your Task 4 sources, which must build with `make -C workloads` in the course container;
- **results/**: one directory per reported run with its raw `stats.txt`, named according to the convention of `scripts/run_all.sh`, i.e., `<cpu>-<workload>[-<variant>]` (the graders spot-check these against your report);
- **reproduce.sh**: regenerates every reported number and plot inside the pinned container without manual steps.

Moodle submission link:

<https://moodle.bilkent.edu.tr/current/mod/assign/view.php?id=21222&forceview=1>

Reproducibility is graded as a gate: if `reproduce.sh` fails to run in the pinned container, your total is capped at 50/100. We will notify you, and you may resubmit a fix before the late-submission cut-off; the late penalty applies to the resubmission.

Grading

Component	Points
Task 1: environment bring-up & stats scavenger hunt	10
Task 2: baseline characterization	20
Task 3: pipeline bring-up, execution time, probe, trace	25
Task 4: hazard microbenchmarks	25
Task 5: report quality & methodology	20
Total	100

Late policy: Late submissions are accepted with a **10% grade penalty per week** (every started week counts as a full week) until the late-submission cut-off for all homeworks: **December 27, 2026, 23:59 (Ankara time)**. No submissions are accepted after the cut-off.

Collaboration policy: This homework is **individual** work. Partners are not allowed; all exercises must be completed and submitted individually. Discussing concepts is fine; sharing code, configurations, measurements, or report text is not.